

H5P2

November 1, 2024

```
[1]: from torchvision import datasets
      from torchvision.transforms import ToTensor
      from torch.utils.data import DataLoader
      import torch
      from torch import nn
      import numpy as np
      import matplotlib.pyplot as plt
```

```
[2]: # Load the USPS dataset
      train_data = datasets.USPS(root='usps', download=True, transform=ToTensor(),
      ↪train=True)
      test_data = datasets.USPS(root='usps', download=True, transform=ToTensor(),
      ↪train=False)

      # Create DataLoaders for training and testing
      train_loader = DataLoader(train_data, batch_size=1024, shuffle=True)
      test_loader = DataLoader(test_data, batch_size=len(test_data), shuffle=False)
```

Downloading

<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multiclass/usps.bz2> to
usps/usps.bz2

100%| | 6.58M/6.58M [00:01<00:00, 5.99MB/s]

Downloading

<https://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/multiclass/usps.t.bz2> to
usps/usps.t.bz2

100%| | 1.83M/1.83M [00:00<00:00, 2.16MB/s]

```
[3]: # Define the MLP model with 1 hidden layer of width 64, following the specified
      ↪structure
      class MLP(nn.Module):
          def __init__(self):
              super().__init__()
              self.flatten = nn.Flatten()
              self.mlp = nn.Sequential(
                  nn.Linear(16 * 16, 64), # Input layer to hidden layer
                  nn.ReLU(),
```

```

        nn.Linear(64, 64),          # Hidden layer
        nn.ReLU(),
        nn.Linear(64, 10)          # Output layer
    )

    def forward(self, X):
        return self.mlp(self.flatten(X))

model = MLP().to('cuda')
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.5)

```

```

[4]: %%time
epochs = 1000
CE = torch.zeros((epochs))
Training = torch.zeros((epochs))
Test = torch.zeros((epochs))
for epoch in range(epochs):
    cumulative_accuracy = 0
    cumulative_loss = 0
    for X, Y in train_loader:
        X, Y = X.to('cuda'), Y.to('cuda')
        out = model(X)
        loss = loss_fn(out, Y)
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        cumulative_loss += loss.item()
        cumulative_accuracy += (out.argmax(axis=1) == Y).sum().item()
    CE[epoch] = cumulative_loss / len(train_loader)
    Training[epoch] = cumulative_accuracy / len(train_data)
    with torch.no_grad():
        for Xt, Yt in test_loader:
            Xt, Yt = Xt.to('cuda'), Yt.to('cuda')
            test_out = model(Xt)
            test_accuracy_epoch = (test_out.argmax(axis=1) == Yt).sum().item() /
↪ len(test_data)
        Test[epoch] = test_accuracy_epoch
    if (epoch + 1) % 100 == 0:
        print(f"Epoch {epoch + 1}/{epochs} | Loss: {CE[epoch]:.4f} | Training_
↪ Accuracy: {Training[epoch]:.4f} | Test Accuracy: {Test[epoch]:.4f}")

```

```

Epoch 100/1000 | Loss: 0.0483 | Training Accuracy: 0.9868 | Test Accuracy:
0.9387
Epoch 200/1000 | Loss: 0.0355 | Training Accuracy: 0.9894 | Test Accuracy:
0.9287
Epoch 300/1000 | Loss: 0.0086 | Training Accuracy: 0.9992 | Test Accuracy:

```

```
0.9412
Epoch 400/1000 | Loss: 0.0047 | Training Accuracy: 0.9996 | Test Accuracy:
0.9417
Epoch 500/1000 | Loss: 0.0031 | Training Accuracy: 0.9999 | Test Accuracy:
0.9422
Epoch 600/1000 | Loss: 0.0023 | Training Accuracy: 0.9999 | Test Accuracy:
0.9432
Epoch 700/1000 | Loss: 0.2746 | Training Accuracy: 0.9172 | Test Accuracy:
0.8475
Epoch 800/1000 | Loss: 0.1600 | Training Accuracy: 0.9510 | Test Accuracy:
0.8909
Epoch 900/1000 | Loss: 0.0972 | Training Accuracy: 0.9639 | Test Accuracy:
0.9088
Epoch 1000/1000 | Loss: 0.0675 | Training Accuracy: 0.9809 | Test Accuracy:
0.9098
CPU times: user 12min 14s, sys: 2.01 s, total: 12min 16s
Wall time: 12min 24s
```

```
[5]: epochs_range = np.arange(1, epochs + 1)
CE_np = CE.cpu().numpy()
Training_np = Training.cpu().numpy()
Test_np = Test.cpu().numpy()

plt.figure(figsize=(10, 5))
plt.plot(epochs_range, CE_np, label="Cross Entropy Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Cross Entropy Loss over Epochs")
plt.legend()
plt.show()

plt.figure(figsize=(10, 5))
plt.plot(epochs_range, Training_np, label="Training Accuracy")
plt.plot(epochs_range, Test_np, label="Test Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Training and Test Accuracy over Epochs")
plt.legend()
plt.show()
```

